

# Efficient statistical learning of complex data

Duzhe Wang


University of Wisconsin-Madison  
Department of Statistics

Ph.D. Thesis Defense  
November 30, 2020

Dissertation committee: Prof. Po-Ling Loh,  
Prof. Varun Jog, Prof. Hyunseung Kang,  
Prof. Vivak Patel, Prof. Anru Zhang

# Thesis overview

## The complex data zoo

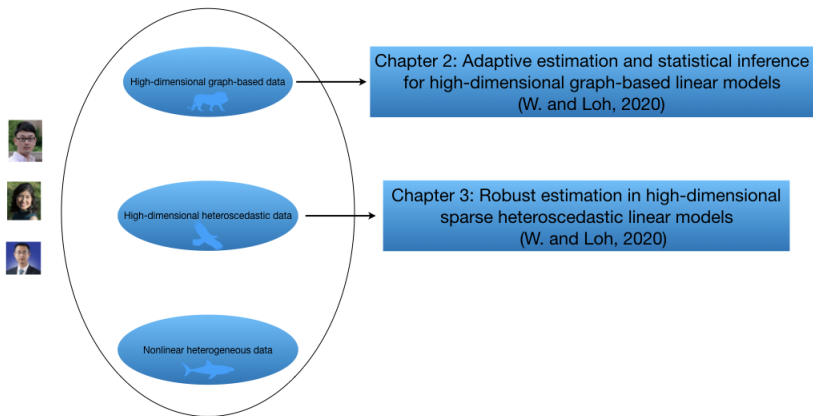


Thanks for showing me around. It's a fun and memorable trip!

# Thesis overview

The complex data zoo

Statistical problems

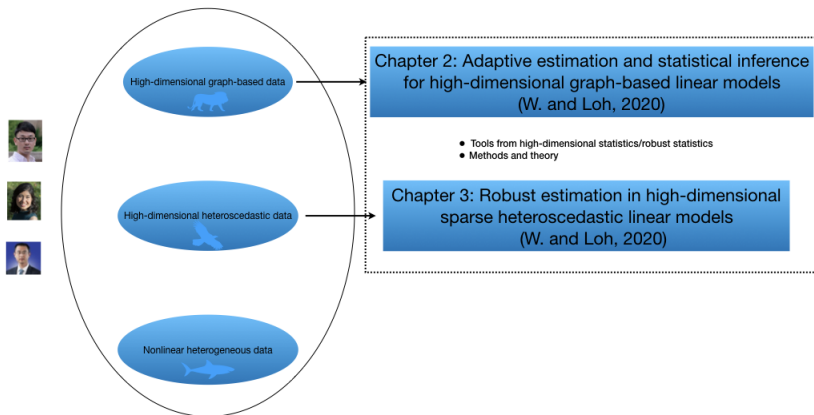


Thanks for showing me around. It's a fun and memorable trip!

# Thesis overview

## The complex data zoo

## Statistical problems

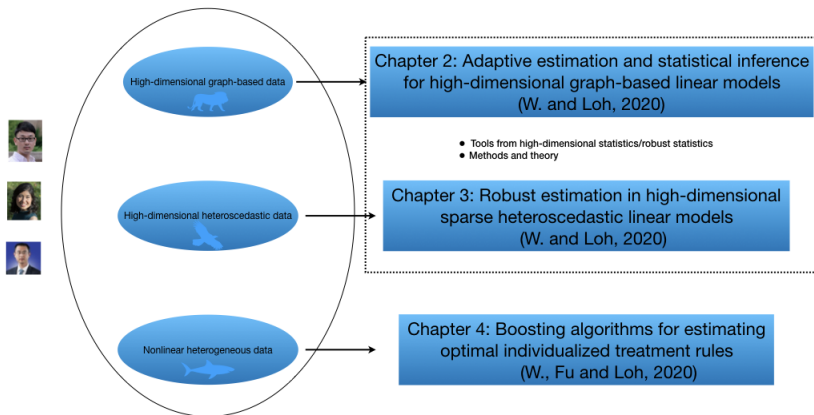


Thanks for showing me around. It's a fun and memorable trip!

# Thesis overview

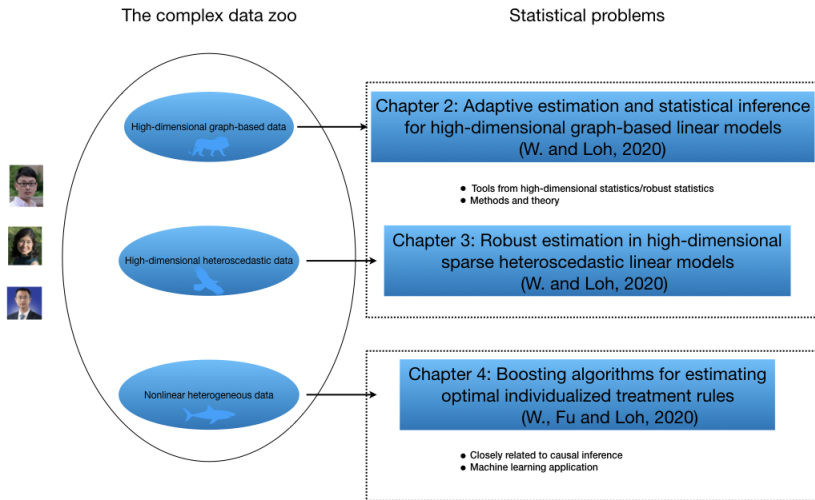
## The complex data zoo

## Statistical problems



Thanks for showing me around. It's a fun and memorable trip!

# Thesis overview



Thanks for showing me around. It's a fun and memorable trip!

# Outline of today's talk

- High-dimensional linear models
- Beyond linear models

# Outline of today's talk

- High-dimensional linear models




- Beyond linear models



# Outline of today's talk


- High-dimensional linear models



- Beyond linear models

# Outline of today's talk


- High-dimensional linear models




- Beyond linear models

# Outline of today's talk

- High-dimensional linear models




- Beyond linear models




# Outline of today's talk

- High-dimensional linear models



- Beyond linear models



# High-dimensional sparse heteroscedastic linear models

$$y = X\beta^* + \varepsilon$$

- $X = (X_1, \dots, X_N)^T \in \mathbb{R}^{N \times n}$ : sub-Gaussian design matrix with  $X_i \in \mathbb{R}^n$
- $y = (y_1, \dots, y_N)^T \in \mathbb{R}^N$ : response vector
- $\beta^* \in \mathbb{R}^n$ : true  $s$ -sparse regression coefficients
- $\varepsilon = (\varepsilon_1, \dots, \varepsilon_N)^T \in \mathbb{R}^N$ :  $\varepsilon_i$ 's are independent, conditionally normal random variables with

$$\mathbb{E}(\varepsilon_i | X_i) = 0, \quad \mathbb{E}(\varepsilon_i^2 | X_i) = W_i$$

- High-dimensional setting:  $N \ll n$
- Conditional heteroscedasticity:  $W_i = g(X_i, \beta^*)$ 
  - Parametric function form of  $g$  is known
  - An example in PET imaging:  $W_i = c|X_i^T \beta^*|$  (Jia, Rohe and Yu, 2013)
- Lasso:

$$\hat{\beta} = \operatorname{argmin}_{\beta \in \mathbb{R}^n} \frac{1}{2N} \|y - X\beta\|_2^2 + \lambda \|\beta\|_1.$$

# Estimation consistency

## Theorem

Assume that  $0 < L_1 \leq W_i \leq L_2 < \infty$  for  $1 \leq i \leq N$ . If  $\lambda \gtrsim \sqrt{\frac{L_2 \log n}{N}}$ , then with high probability, we have  $\|\hat{\beta} - \beta^*\|_2 \lesssim \lambda \sqrt{s}$  and  $\|\hat{\beta} - \beta^*\|_1 \lesssim \lambda s$ .

# Estimation consistency

## Theorem

Assume that  $0 < L_1 \leq W_i \leq L_2 < \infty$  for  $1 \leq i \leq N$ . If  $\lambda \gtrsim \sqrt{\frac{L_2 \log n}{N}}$ , then with high probability, we have  $\|\hat{\beta} - \beta^*\|_2 \lesssim \lambda \sqrt{s}$  and  $\|\hat{\beta} - \beta^*\|_1 \lesssim \lambda s$ .

- If  $\lambda \asymp \sqrt{\frac{L_2 \log n}{N}}$ , then

$$\|\hat{\beta} - \beta^*\|_2 \leq \mathcal{O}_{\mathbb{P}} \left( \sqrt{\frac{s L_2 \log n}{N}} \right) \quad \text{and} \quad \|\hat{\beta} - \beta^*\|_1 \leq \mathcal{O}_{\mathbb{P}} \left( s \sqrt{\frac{L_2 \log n}{N}} \right)$$

- The Lasso is  $\ell_2$ -consistent if  $\frac{s L_2 \log n}{N} = o(1)$
- If  $W_i = \sigma_\varepsilon^2$  for  $i = 1, \dots, N$  (homoscedastic case), then

$$\|\hat{\beta} - \beta^*\|_2 \leq \mathcal{O}_{\mathbb{P}} \left( \sigma_\varepsilon \sqrt{\frac{s \log n}{N}} \right) \quad \text{and} \quad \|\hat{\beta} - \beta^*\|_1 \leq \mathcal{O}_{\mathbb{P}} \left( \sigma_\varepsilon s \sqrt{\frac{\log n}{N}} \right)$$

# Variable selection consistency

## Mutual incoherence condition

Let  $S$  denote the support set of  $\beta^*$ . For covariance matrix  $\Sigma_x$ , there exists a constant  $\alpha \in (0, 1)$ , such that  $\|(\Sigma_x)_{S^c S}(\Sigma_x)_{SS}^{-1}\|_\infty \leq \frac{\alpha}{2}$ .

## Theorem

Assume that mutual incoherence condition holds. Let  $\widehat{S}$  be the support set of  $\widehat{\beta}$ . Let  $\lambda \gtrsim \frac{4}{1-\alpha} \sqrt{\frac{L_2 \log(n-s)}{N}}$ .

- (a) With high probability, we have  $\widehat{S} \subseteq S$ .
- (b) If  $\min_{i \in S} |\beta_i^*| > \sqrt{\frac{8}{\lambda_{\min}(\Sigma_x)}} \sqrt{\frac{L_2 \log s}{N}} + \lambda (\|(\Sigma_x)_{SS}^{-1}\|_\infty + cs\sqrt{\frac{s}{N}})$ , then with high probability, we have  $S = \widehat{S}$ .

---

$$\|M\|_\infty = \max_{i=1, \dots, m} \sum_{j=1}^n |M_{ij}|$$



(1) Characterizing (asymptotic) distribution of the Lasso is difficult

- KKT condition:

$$M\sqrt{N}(\hat{\beta} - \beta^*) + \sqrt{N}\lambda\hat{k} = \frac{1}{\sqrt{N}}X^T\varepsilon,$$

where  $M = \frac{1}{N}X^TX$  and  $\hat{k}$  be the subgradient of  $\|\cdot\|_1$  at  $\hat{\beta}$

- Low-dimensional setting:


$$\sqrt{N}(\hat{\beta} - \beta^*) = \frac{1}{\sqrt{N}}M^{-1}X^T\varepsilon - \sqrt{N}\lambda M^{-1}\hat{k}$$

- High-dimensional setting: hard to characterize  $\sqrt{N}(\hat{\beta} - \beta^*)$


(2) No side/prior information on heteroscedasticity is used in the Lasso

# Outline of today's talk

- High-dimensional linear models



- Beyond linear models



# Robustifying the Lasso under heteroscedasticity

## Algorithm

Input: dataset  $\{(X_i, y_i)\}_{i=1}^N$ , formula of  $g$ , tuning parameters  $\lambda$  and  $\mu$ .

- (1) Solve the Lasso to obtain a preliminary estimator  $\hat{\beta}$ .
- (2) Set  $\widehat{W}_i = g(X_i, \hat{\beta})$ ,  $\widehat{W} = \text{diag}(\widehat{W}_1, \dots, \widehat{W}_N)$ , and  $\widehat{\Sigma}_N = \frac{1}{N} X^T \widehat{W}^{-1} X$ .
- (3) Solve the following optimization problem to obtain  $\widehat{\Theta}$ :

$$\begin{aligned} \widehat{\Theta} = \operatorname{argmin}_{\Theta \in \mathbb{R}^{n \times n}} \quad & \|\Theta\|_{1,1} = \sum_{i=1}^n \sum_{j=1}^n |\Theta_{ij}| \\ \text{subject to} \quad & \|\Theta \widehat{\Sigma}_N - I_n\|_{\infty} \leq \mu. \end{aligned}$$

- (4) Output the final estimator:

$$\tilde{\beta} = \hat{\beta} + \frac{1}{N} \widehat{\Theta} X^T \widehat{W}^{-1} (y - X \hat{\beta}).$$

# Robustifying the Lasso under heteroscedasticity

## Algorithm

Input: dataset  $\{(X_i, y_i)\}_{i=1}^N$ , formula of  $g$ , tuning parameters  $\lambda$  and  $\mu$ .

- (1) Solve the Lasso to obtain a preliminary estimator  $\hat{\beta}$ .
- (2) Set  $\widehat{W}_i = g(X_i, \hat{\beta})$ ,  $\widehat{W} = \text{diag}(\widehat{W}_1, \dots, \widehat{W}_N)$ , and  $\widehat{\Sigma}_N = \frac{1}{N} X^T \widehat{W}^{-1} X$ .
- (3) Solve the following optimization problem to obtain  $\widehat{\Theta}$ :

$$\begin{aligned} \widehat{\Theta} = \underset{\Theta \in \mathbb{R}^{n \times n}}{\text{argmin}} \quad & \|\Theta\|_{1,1} = \sum_{i=1}^n \sum_{j=1}^n |\Theta_{ij}| \\ \text{subject to} \quad & \|\Theta \widehat{\Sigma}_N - I_n\|_{\infty} \leq \mu. \end{aligned}$$

$$\begin{aligned} \widehat{\Theta} &= (\hat{\theta}_1, \dots, \hat{\theta}_n)^T \\ \hat{\theta}_i &= \underset{\theta \in \mathbb{R}^n}{\text{argmin}} \quad \|\theta\|_1 \\ \text{subject to} \quad & \|\widehat{\Sigma}_N \theta - e_i\|_{\infty} \leq \mu. \end{aligned}$$

- (4) Output the final estimator:

$$\tilde{\beta} = \hat{\beta} + \frac{1}{N} \widehat{\Theta} X^T \widehat{W}^{-1} (y - X \hat{\beta}).$$

- One-step MLE in the low-dimensional setting:

$$\hat{\beta}_1 = \hat{\beta}_0 - [\nabla S(\hat{\beta}_0)]^{-1} S(\hat{\beta}_0),$$

where  $S(\beta)$  is the score function and  $\hat{\beta}_0$  is an initial estimator of  $\beta^*$

- If  $W$  is known, then in the low-dimensional setting, one-step MLE is

$$\hat{\beta}_1 = \hat{\beta}_0 + \left( \frac{1}{N} X^T W^{-1} X \right)^{-1} \frac{1}{N} X^T W^{-1} (y - X \hat{\beta}_0)$$

- One-step MLE in the low-dimensional setting:

$$\hat{\beta}_1 = \hat{\beta}_0 - [\nabla S(\hat{\beta}_0)]^{-1} S(\hat{\beta}_0),$$

where  $S(\beta)$  is the score function and  $\hat{\beta}_0$  is an initial estimator of  $\beta^*$

- If  $W$  is known, then in the low-dimensional setting, one-step MLE is

$$\hat{\beta}_1 = \hat{\beta}_0 + \left( \frac{1}{N} X^T W^{-1} X \right)^{-1} \frac{1}{N} X^T W^{-1} (y - X \hat{\beta}_0)$$

- Challenges:

- (1) In the high-dimensional setting,  $\frac{1}{N} X^T W^{-1} X$  is singular
- (2) In the heteroscedastic setting,  $W$  is unknown

- Solutions:

- (1) Find a sparse approximate inverse
- (2) Utilize the side information to estimate  $W$

## Theorem

Assume that  $\Sigma^* = \mathbb{E} \left( \frac{X_i X_i^T}{W_i} \right)$  is positive definite. Let  $\Theta^* \in \mathbb{R}^{n \times n}$  denote the inverse of  $\Sigma^*$ . Let

$$\lambda \asymp \sqrt{\frac{L_2 \log n}{N}} \quad \text{and} \quad \mu \asymp \frac{1}{L_1} \sqrt{\frac{\log n}{N}} + s \sqrt{\frac{L_2 \log n}{N}}.$$

Under a set of assumptions, we have for  $1 \leq j \leq n$ ,

$$\sqrt{N} (\tilde{\beta}_j - \beta_j^*) \xrightarrow{d} N(0, e_j^T \Theta^* e_j).$$

- Similar arguments yield

$$\frac{\sqrt{N}(\tilde{\beta}_j - \beta_j^*)}{\sqrt{e_j^T \widehat{\Theta} \widehat{\Sigma}_N \widehat{\Theta}^T e_j}} \rightarrow N(0, 1)$$

- Confidence interval: let  $\Phi(x)$  be the cumulative distribution function of  $N(0, 1)$ . Then

$$\left[ \tilde{\beta}_j - \Phi^{-1}\left(1 - \frac{\alpha}{2}\right) \sqrt{\frac{e_j^T \widehat{\Theta} \widehat{\Sigma}_N \widehat{\Theta}^T e_j}{N}}, \tilde{\beta}_j + \Phi^{-1}\left(1 - \frac{\alpha}{2}\right) \sqrt{\frac{e_j^T \widehat{\Theta} \widehat{\Sigma}_N \widehat{\Theta}^T e_j}{N}} \right]$$

provides an asymptotically valid  $(1 - \alpha)$ -confidence interval for  $\beta_j^*$

- Hypothesis testing: test statistic for testing whether  $\beta_j^*$  is equal to 0






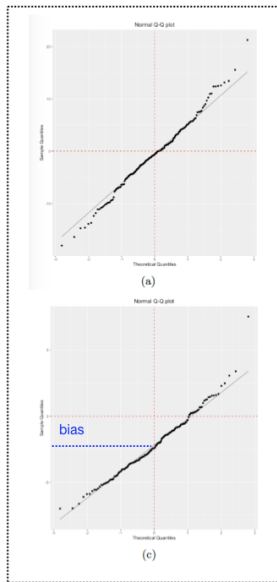
# Simulation settings

- $(N, n) = (120, 150)$
- $\beta^* = (3, 4, 3, 1.5, 2, 1.5, 0, \dots, 0)^T$
- $X_i \sim N(0, \Sigma_x)$ , where  $(\Sigma_x)_{S^c S} = 0$ , and  $(\Sigma_x)_{ij} = 0.5^{|i-j|}$  if both  $i$  and  $j$  are in  $S$  or both  $i$  and  $j$  are in  $S^c$
- $\varepsilon_i \mid X_i \sim N(0, W_i)$ , where

$$W_i = \min\left(\frac{1}{25} \exp\left(\frac{1}{2} |X_i^T \beta^*|\right), 5\right)$$

- $y_i = X_i^T \beta^* + \varepsilon_i$

# Simulation results





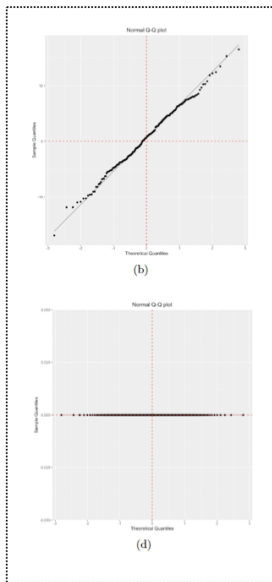

(a) Q-Q plot of  $\sqrt{N} (\tilde{\beta}_5 - \beta_5^*)$

(b) Q-Q plot of  $\sqrt{N} (\tilde{\beta}_7 - \beta_7^*)$

(c) Q-Q plot of  $\sqrt{N} (\hat{\beta}_5 - \beta_5^*)$

(d) Q-Q plot of  $\sqrt{N} (\hat{\beta}_7 - \beta_7^*)$

# Simulation results



(a) Q-Q plot of  $\sqrt{N} (\tilde{\beta}_5 - \beta_5^*)$


(b) Q-Q plot of  $\sqrt{N} (\tilde{\beta}_7 - \beta_7^*)$

(c) Q-Q plot of  $\sqrt{N} (\hat{\beta}_5 - \beta_5^*)$


(d) Q-Q plot of  $\sqrt{N} (\hat{\beta}_7 - \beta_7^*)$

# Outline of today's talk

- High-dimensional linear models



- Beyond linear models





# Motivating examples

- Gene expression



Metabolic pathways (cited from Wikipedia)

- Identify genes which are associated with a target gene among hundreds of candidates from a large number of metabolic pathways
- Genes within a same cluster (pathway) have similar patterns
- Model:  $y = X\beta + \epsilon$ 
  - $y$ : expression levels of the target gene
  - $X$ : design matrix including expression levels of candidate genes
  - $\beta$ : association levels between genes

# Motivating examples

- Gene expression




Metabolic pathways (cited from Wikipedia)

- Identify genes which are associated with a target gene among hundreds of candidates from a large number of metabolic pathways
- Genes within a same cluster (pathway) have similar patterns
- Model:  $y = X\beta + \epsilon$ 
  - $y$ : expression levels of the target gene
  - $X$ : design matrix including expression levels of candidate genes
  - $\beta$ : association levels between genes

- Conclusion: these high-dimensional datasets have structures that can be captured in the form of complex graphs

# High-dimensional graph-based linear models

$$y = X\beta^* + \varepsilon$$




- $X = (X_1, \dots, X_N)^T \in \mathbb{R}^{N \times n}$ : design matrix with  $X_i \in \mathbb{R}^n$
- $y = (y_1, \dots, y_N)^T \in \mathbb{R}^N$ : response vector
- $\beta^* \in \mathbb{R}^n$ : unknown true regression coefficients
- $\varepsilon = (\varepsilon_1, \dots, \varepsilon_N)^T \in \mathbb{R}^N$ : i.i.d. random error
- High-dimensional setting:  $N \ll n$
- Graph setting: coordinates of  $\beta^*$  correspond to nodes of some known underlying undirected graph
- Goal: estimate sparse regression coefficients with certain graph-based structure



# Graph difference operator

Let  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$  be an undirected graph with  $|\mathcal{V}| = n$  and  $|\mathcal{E}| = p$ .

- Oriented incidence matrix  $F \in \{-1, 0, 1\}^{p \times n}$ : if the  $k$ th edge is  $(i, j) \in \mathcal{E}$  with  $i < j$ , then the  $k$ th row of  $F$  is  $(0, \dots, -1, \dots, +1, \dots, 0)$



$$F = \begin{bmatrix} -1 & 1 & 0 & 0 & 0 & 0 \\ 0 & -1 & 1 & 0 & 0 & 0 \\ 0 & -1 & 0 & 1 & 0 & 0 \\ 0 & 0 & -1 & 1 & 0 & 0 \\ 0 & 0 & -1 & 0 & 1 & 0 \\ 0 & 0 & 0 & -1 & 1 & 0 \\ 0 & 0 & 0 & 0 & -1 & 1 \end{bmatrix}$$

- Laplacian matrix:  $L = F^T F \in \mathbb{R}^{n \times n}$
- Higher-order graph difference operator: graph difference operator of order  $k + 1$  is

$$\Delta^{(k+1)} = \begin{cases} F^T \Delta^{(k)} = L^{\frac{k+1}{2}} \in \mathbb{R}^{n \times n} & \text{for odd } k \\ F \Delta^{(k)} = FL^{\frac{k}{2}} \in \mathbb{R}^{p \times n} & \text{for even } k. \end{cases}$$

# Graph-based structure

## Graph-based piecewise polynomial structure (Wang et al., 2016)

for  $k \geq 0$  and  $s > 0$ ,  $\beta^*$  is  $(k, s)$ -piecewise polynomial over the graph  $\mathcal{G}$  if

$$\|\Delta^{(k+1)}\beta^*\|_0 \leq s$$

- $k = 0$ : piecewise constant
- $k = 1$ : piecewise linear
- $k = 2$ : piecewise quadratic
- Local smoothness
- Focus on simultaneously sparse and piecewise polynomial regression coefficients:

$$\beta^* \in \mathcal{S}(k, s_1, s_2) = \left\{ \beta \in \mathbb{R}^n : \|\Delta^{(k+1)}\beta\|_0 \leq s_1, \|\beta\|_0 \leq s_2 \right\}$$

## $k$ -th order Graph-Piecewise-Polynomial-Lasso

$$\hat{\beta}_g = \operatorname{argmin}_{\beta \in \mathbb{R}^n} \frac{1}{2N} \|y - X\beta\|_2^2 + \lambda \|\beta\|_1 + \lambda_g \|\Delta^{(k+1)}\beta\|_1$$


- $k = 0$ : Graph-Fused Lasso (Kim, Sohn and Xing, 2019)
- Assume  $k$  is known for theory, but tune  $k$  in practice
- Equivalent formulation:

$$\hat{\beta}_g = \operatorname{argmin}_{\beta \in \mathbb{R}^n} \frac{1}{2N} \|y - X\beta\|_2^2 + \lambda \|D\beta\|_1,$$

$$\text{where } D = \begin{bmatrix} \frac{\lambda_g}{\lambda} \Delta^{(k+1)} \\ I_n \end{bmatrix}$$

- Graph-Piecewise-Polynomial-Lasso can be solved efficiently via the ADMM algorithm

# Simulation settings



- Coordinates of  $\beta^*$  correspond to a 2d grid graph with 25 rows and 25 columns ( $n = 625, p = 1200$ )
- $\beta^* \in \mathbb{R}^{625}$  is sparse and piecewise constant ( $s_1 = 54, s_2 = 81$ )
- Construction of  $\beta^*$ : first construct  $B^* \in \mathbb{R}^{25 \times 25}$ , then stack columns of  $B^*$  on top of one another
- $X_i \sim N(0, I_{n \times n}); \varepsilon_i \sim N(0, 0.1); y_i = X_i^T \beta^* + \varepsilon_i$
- $N = 250$


# Simulation results

- Graph-Smooth-Lasso:


$$\hat{\beta}^{\text{gsmooth}} = \underset{\beta \in \mathbb{R}^n}{\operatorname{argmin}} \frac{1}{2N} \|y - X\beta\|_2^2 + \lambda_1 \|\beta\|_1 + \lambda_2 \|\Delta^{(1)}\beta\|_2^2$$

- Graph-Spline-Lasso:


$$\hat{\beta}^{\text{gspline}} = \underset{\beta \in \mathbb{R}^n}{\operatorname{argmin}} \frac{1}{2N} \|y - X\beta\|_2^2 + \lambda_1 \|\beta\|_1 + \lambda_2 \|\Delta^{(2)}\beta\|_2^2$$




Our approach



Lasso




Graph-Smooth-Lasso




Graph-Spline-Lasso

# Outline of today's talk

- High-dimensional linear models



- Beyond linear models




# A motivating example

- COVID-19 patients are a very heterogeneous population

## Death rates depend on age group

Estimated proportion of deaths among infections




Source: Imperial College London, 16 March, SAGE

BBC

## Death rates depend on underlying health

Proportion of deaths among confirmed cases



Source: Chinese Centre for Disease Control and Prevention, Feb 18

BBC

# A motivating example

- Heterogeneous treatment effects:

*“Primary efficacy analysis demonstrates BNT162b2 to be 95% effective against COVID-19 beginning 28 days after the first dose; 170 confirmed cases of COVID-19 were evaluated, with 162 observed in the placebo group versus 8 in the vaccine group”*


—Phase 3 study results of Pfizer and BioNTech's vaccine

*“Today's primary analysis was based on 196 cases, of which 185 cases of COVID-19 were observed in the placebo group versus 11 cases observed in the mRNA-1273 group, resulting in a point estimate of vaccine efficacy of 94.1%.”*

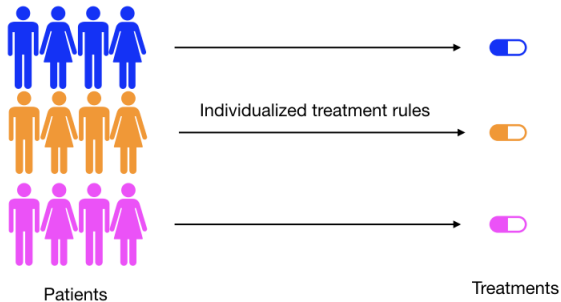
—Phase 3 study results of Moderna's vaccine



# Hope of precision medicine



# Hope of precision medicine



# Individualized treatment rules

- $\{(X_i, A_i, Y_i), 1 \leq i \leq n\}$ : i.i.d. observations of  $(X, A, Y)$ 
  - $X \subset \mathcal{X} \subset \mathbb{R}^p$ : prognostic variables
  - $A \subset \mathcal{A} = \{-1, +1\}$ : the given treatment
  - $Y \subset \mathbb{R}$ : the patient clinical outcome (with larger being better)
- Individualized treatment rule (ITR):

$$\mathcal{D} : \mathcal{X} \rightarrow \{-1, +1\}$$

- e.g.,  $\mathcal{D}(x) = 1, \mathcal{D}(x) = \text{sign}(x^T \mathbf{1})$

# Individualized treatment rules

- $\{(X_i, A_i, Y_i), 1 \leq i \leq n\}$ : i.i.d. observations of  $(X, A, Y)$ 
  - $X \subset \mathcal{X} \subset \mathbb{R}^p$ : prognostic variables
  - $A \subset \mathcal{A} = \{-1, +1\}$ : the given treatment
  - $Y \subset \mathbb{R}$ : the patient clinical outcome (with larger being better)
- Individualized treatment rule (ITR):

$$\mathcal{D} : \mathcal{X} \rightarrow \{-1, +1\}$$

- e.g.,  $\mathcal{D}(x) = 1, \mathcal{D}(x) = \text{sign}(x^T \mathbf{1})$
- The optimal ITR  $\mathcal{D}^*(x)$ :

$$\mathcal{D}^*(x) = \operatorname{argmax}_{a \in \mathcal{A}} \underbrace{Q(x, a) := \mathbb{E}(Y|x, a)}_{\text{Quality}}$$

# Two learning frameworks

## Generic method of indirect learning

- (1) Assume  $Q(x, 1)$  and  $Q(x, -1)$  are in some specified functional space  $\mathcal{F}$
- (2) Estimate  $Q(x, 1)$  and  $Q(x, -1)$ : this is a regression problem
- (3) Estimated optimal ITR:

$$\widehat{\mathcal{D}}(x) = \text{sign}(\widehat{Q}(x, 1) - \widehat{Q}(x, -1))$$

## Generic method of direct learning


- (1) Note  $\mathcal{D}^*(x) = \text{sign}(f^*(x))$ . Assume  $f^*(x) \in \mathcal{F}$
- (2) Estimate  $f^*(x)$ : this is a classification problem
- (3) Estimated optimal ITR:

$$\widehat{\mathcal{D}}(x) = \text{sign}(\hat{f}(x))$$

- Goal: estimate the nonlinear and complex optimal ITR  $\mathcal{D}^*(x)$  with the observed dataset

# Outline of today's talk

- High-dimensional linear models



- Beyond linear models



# Our first proposed method

- One instance of indirect learning
- Key ideas:
  - Additive regression trees: assume

$$Q(x, 1) = \sum_{t=1}^K b_1^{(t)}(x),$$

and

$$Q(x, -1) = \sum_{t=1}^K b_{-1}^{(t)}(x),$$

where  $b_1^{(t)}(x)$  and  $b_{-1}^{(t)}(x)$  are regression trees

- Use boosting algorithm to estimate regression trees sequentially

# XGBoost algorithm

- 1st iteration:

## Estimation of $b_1^{(1)}$

- (1) Fit a tree to the training data  $(X_i, Y_i)$ :

$$\hat{f}^{(1)} = \underset{f}{\operatorname{argmin}} \sum_{i:A_i=1} (Y_i - f(X_i))^2 + J(f)$$

- $f(x)$  is a regression tree:  $f(x) = w_{q(x)}(q: \mathbb{R}^p \rightarrow T, w \in \mathbb{R}^{|T|})$ , where  $q$  represents the tree structure and  $T$  represents the leaves
- $J(f)$  is the cost complexity of a regression tree:  $J(f) = \gamma|T| + \frac{1}{2}\lambda\|w\|_2^2$

- (2) Shrinkage:  $\hat{b}_1^{(1)} = \eta \hat{f}^{(1)}$ , where  $0 < \eta < 1$



# XGBoost algorithm

- $t$ th iteration ( $t > 1$ ):

## Estimation of $b_1^{(t)}$

- (1) After  $(t - 1)$ th iteration:  $\hat{Y}_i^{(t-1)} = \sum_{k=1}^{t-1} \hat{b}_1^{(k)}(X_i)$  is the estimated outcome value of  $X_i$
- (2) Fit a tree to the training data  $(X_i, Y_i)$ :


$$\hat{f}^{(t)} = \underset{f}{\operatorname{argmin}} \sum_{i:A_i=1} [Y_i - (\hat{Y}_i^{(t-1)} + f(X_i))]^2 + J(f),$$

- (3) Shrinkage:  $\hat{b}_1^{(t)} = \eta \hat{f}^{(t)}$

- Output the boosted model:

$$\widehat{Q}(x, 1) = \sum_{t=1}^K \hat{b}_1^{(t)}(x)$$


# How do we fit a regression tree?



- Decide optimal leaf weights: for a fixed tree structure  $T$ , let  $I_j = \{i | q(X_i) = j\}$  be the instance set of leaf  $j$ . Then

$$w_j^* = \frac{2 \sum_{i \in I_j} (Y_i - \hat{Y}_i^{(t-1)})}{2|I_j| + \lambda}$$

# How do we fit a regression tree?



- Decide optimal leaf weights: for a fixed tree structure  $T$ , let  $I_j = \{i | q(X_i) = j\}$  be the instance set of leaf  $j$ . Then

$$w_j^* = \frac{2 \sum_{i \in I_j} (Y_i - \hat{Y}_i^{(t-1)})}{2|I_j| + \lambda}$$

- Split finding algorithm for estimating tree structure  $T$ :  
Chen and Guestrin, 2016

## Algorithm

Input: dataset  $\{(X_i, Y_i, A_i)\}_{i=1}^n$ , number of iterations  $K$ , learning rate  $\eta$ , maximum of tree depth  $d$

- (1) Train  $\text{bst.plus1} = \text{XGBoost}(\{(X_i, Y_i); A_i = 1\}, K, \eta, d)$
- (2) Train  $\text{bst.minus1} = \text{XGBoost}(\{(X_i, Y_i); A_i = -1\}, K, \eta, d)$
- (3) Output the estimated optimal ITR:

$$\widehat{\mathcal{D}}(x) = \text{sign}(\text{bst.plus1}(x) - \text{bst.minus1}(x))$$

## Our second proposed method

- One instance of direct learning
- Key ideas:
  - Assume  $f^*(x) = \sum_{t=1}^K b^{(t)}(x)$  where  $b^{(t)}$  are regression trees
  - Use boosting algorithm to estimate  $b^{(t)}$  sequentially

# Our second proposed method

- One instance of direct learning
- Key ideas:
  - Assume  $f^*(x) = \sum_{t=1}^K b^{(t)}(x)$  where  $b^{(t)}$  are regression trees
  - Use boosting algorithm to estimate  $b^{(t)}$  sequentially

## Fisher consistency theorem

Assume  $Y = \mu(X) + \delta(X) \times A + \varepsilon$ . Let  $\pi_A(X) = P(A|X)$ . Then we have

$$\mu = \operatorname{argmin}_g \mathbb{E} \left\{ \frac{1}{\pi_A(X)} (Y - g(X))^2 \right\}.$$

Furthermore,

$$f^* = \operatorname{argmin}_f \mathbb{E} \left\{ \underbrace{\frac{|Y - \mu(X)|}{\pi_A(X)}}_{\text{weight}} \phi \left( \underbrace{A \times \operatorname{sign}(Y - \mu(X))}_{\text{adjusted label}} f(X) \right) \right\},$$

functional margin  
adjusted label

where  $\phi(x) = \log(1 + e^{-2x})$ .

## Estimation of $\mu(x)$

(1) Assume  $\mu(x) = \alpha_0 + \alpha^T x$

(2) Estimate  $\alpha_0$  and  $\alpha$ :

$$\mu = \operatorname{argmin}_g E \left\{ \frac{1}{\pi_A(X)} (Y - g(X))^2 \right\}$$

$$\hat{\alpha}_0, \hat{\alpha} = \operatorname{argmin}_{\alpha_0, \alpha} \sum_{i=1}^n \frac{1}{\pi_{A_i}(X_i)} (Y_i - \alpha_0 - \alpha^T X_i)^2$$

(3) Estimate  $\mu(x)$ :

$$\hat{\mu}(x) = \hat{\alpha}_0 + \hat{\alpha}^T x$$

# XGBoost algorithm

- 1st iteration:

## Estimation of $b^{(1)}$

- (1) Fit a tree to the training data  $(X_i, A_i, Y_i)$ :

$$f^* = \operatorname{argmin}_f \mathbb{E} \left\{ \frac{|Y - \mu(X)|}{\pi_A(X)} \phi(A \times \operatorname{sign}(Y - \mu(X))f(X)) \right\}$$

$$\hat{f}^{(1)} = \operatorname{argmin}_f \sum_{i=1}^n \frac{|Y_i - \hat{\mu}(X_i)|}{\pi_{A_i}(X_i)} \underbrace{\phi(A_i f(X_i) \times \operatorname{sign}(Y_i - \hat{\mu}(X_i)))}_{\text{Use second-order approximation}} + J(f)$$

- (2) Shrinkage:  $\hat{b}_1^{(1)} = \eta \hat{f}^{(1)}$ , where  $0 < \eta < 1$



# XGBoost algorithm

- $t$ th iteration:

## Estimation of $b^{(t)}$

- (1) Fit a tree to the training data  $(X_i, A_i, Y_i)$ :

$$\hat{f}^{(t)} = \underset{f}{\operatorname{argmin}} \sum_{i=1}^n \frac{|Y_i - \hat{\mu}(X_i)|}{\pi_{A_i}(X_i)} \underbrace{\phi \left( A_i \left( \hat{Y}_i^{(t-1)} + f(X_i) \right) \times \operatorname{sign}(Y_i - \hat{\mu}(X_i)) \right)}_{\text{Use second-order approximation}} + J(f)$$

- (2) Shrinkage:  $\hat{b}^{(t)} = \eta \hat{f}^{(t)}$

## Algorithm

Input: dataset  $\{(X_i, A_i, Y_i)\}_{i=1}^n$ , number of iterations  $K$ , shrinkage parameter  $\eta$  and maximum tree depth  $d$ .

- (1) Estimate the common effect  $\mu$ .
- (2) Train  $\text{bst} = \text{XGBoost}(\{X_i, A_i \text{sign}(Y_i - \hat{\mu}(X_i))\}, K, \eta, d)$  with weighted deviance loss
- (3) Output the estimated optimal ITR:

$$\hat{\mathcal{D}}(x) = \text{sign}(\text{bst}(x))$$

# Simulation settings

- $X_i \in \mathbb{R}^{10}$ : each component is i.i.d.  $U(-1, 1)$
- $A_i$ :  $P(A_i = -1) = P(A_i = 1) = 0.5$
- $\varepsilon_i \sim N(0, 1)$
- $Y_i = 1 + 2X_{1i} + X_{2i} + 0.5X_{3i} + \delta(X_i) \times A_i + \varepsilon_i$ , where  $X_{1i}$ ,  $X_{2i}$  and  $X_{3i}$  are the first, second and third components of  $X_i$ , and

$$\delta(X_i) = 0.2 + X_{1i}^2 + X_{2i}^2 - X_{3i}^2 - X_{4i}^2$$

- Polynomial-type optimal ITR:


$$\mathcal{D}^*(X) = \begin{cases} 1 & 0.2 + X_1^2 + X_2^2 - X_3^2 - X_4^2 > 0 \\ -1 & 0.2 + X_1^2 + X_2^2 - X_3^2 - X_4^2 < 0 \end{cases}$$

## Simulation results

- Misclassification rate of an ITR:  $\frac{1}{n} \sum_{i=1}^n I(\mathcal{D}^*(X_i) \neq \mathcal{D}(X_i))$  for a testing dataset  $\{(X_i, A_i, Y_i), 1 \leq i \leq n\}$

# Simulation results

- Misclassification rate of an ITR:  $\frac{1}{n} \sum_{i=1}^n I(\mathcal{D}^*(X_i) \neq \mathcal{D}(X_i))$  for a testing dataset  $\{(X_i, A_i, Y_i), 1 \leq i \leq n\}$




Linear competitors: D-learning (Qi et al., 2019),  $\ell_1$ -PLS (Qian and Murphy, 2011), OWL-Linear (Zhao et al., 2012), Q-learning (benchmark); Nonlinear competitor: OWL-RBF (Zhao et al., 2012)

# Simulation results

- Value function of an ITR:  $V(\mathcal{D}) = \mathbb{E}^{\mathcal{D}}(Y) = \mathbb{E} \left\{ Y \frac{I(A=\mathcal{D}(X))}{\pi_A(X)} \right\}$
- Estimated value function of an ITR:  $\hat{V}(\mathcal{D}) = \frac{\frac{1}{n} \sum_{i=1}^n \frac{Y_i}{\pi_{A_i}(X_i)} I(\mathcal{D}(X_i)=A_i)}{\frac{1}{n} \sum_{i=1}^n \frac{I(\mathcal{D}(X_i)=A_i)}{\pi_{A_i}(X_i)}}$

# Simulation results

- Value function of an ITR:  $V(\mathcal{D}) = \mathbb{E}^{\mathcal{D}}(Y) = \mathbb{E} \left\{ Y \frac{I(A=\mathcal{D}(X))}{\pi_A(X)} \right\}$
- Estimated value function of an ITR:  $\hat{V}(\mathcal{D}) = \frac{\frac{1}{n} \sum_{i=1}^n \frac{Y_i}{\pi_{A_i}(X_i)} I(\mathcal{D}(X_i)=A_i)}{\frac{1}{n} \sum_{i=1}^n \frac{I(\mathcal{D}(X_i)=A_i)}{\pi_{A_i}(X_i)}}$



# Diabetes data analysis

- The dataset was collected from a randomized, double-blind, parallel-group Phase III trial (Charbonnel and Matthews et al., 2005)
- $\mathcal{A} = \{\text{gliclazide, pioglitazone}\}$
- Among 1247 patients, 624 patients received gliclazide and 623 received pioglitazone
- $X$ : 21 pretreatment covariates, e.g., BMI and blood pressure
- $Y$ : primary efficacy endpoint, i.e., change of HbA1c level during 52 weeks




# Diabetes data analysis

- The dataset was collected from a randomized, double-blind, parallel-group Phase III trial (Charbonnel and Matthews et al., 2005)
- $\mathcal{A} = \{\text{gliclazide, pioglitazone}\}$
- Among 1247 patients, 624 patients received gliclazide and 623 received pioglitazone
- $X$ : 21 pretreatment covariates, e.g., BMI and blood pressure
- $Y$ : primary efficacy endpoint, i.e., change of HbA1c level during 52 weeks
- Results:

Method	Our first proposed method	Our second proposed method	Q-learning	I1-PLS	D-learning	OWL-Linear	OWL-RBF
Estimated value	1.447	1.448	1.369	1.428	1.416	1.360	1.363

# Summary

## High-dimensional linear models



### Efficient statistical learning of complex data

#### One-step estimator

Input: dataset  $\{(X_i, y_i)\}_{i=1}^N$ , formula of  $g$ , tuning parameters  $\lambda$  and  $\mu$ .

- (1) Solve the Lasso to obtain a preliminary estimator  $\hat{\beta}$ .
- (2) Set  $\tilde{W}_i = g(X_i, \hat{\beta})$ ,  $\tilde{W} = \text{diag}(\tilde{W}_1, \dots, \tilde{W}_N)$ , and  $\tilde{\Sigma}_N = \frac{1}{N} X^T \tilde{W}^{-1} X$ .
- (3) Solve the following optimization problem to obtain  $\tilde{\Theta}$ :

$$\tilde{\Theta} = \underset{\Theta \in \mathbb{R}^{m \times p}}{\text{argmin}} \quad |\Theta|_{1,1} = \sum_{i=1}^n \sum_{j=1}^p |\Theta_{ij}|$$

subject to  $|\Theta \tilde{\Sigma}_N - I_n|_\infty \leq \mu$ .


- (4) Output the final estimator:

$$\hat{\beta} = \hat{\beta} + \frac{1}{N} \tilde{\Theta} X^T \tilde{W}^{-1} (y - X \hat{\beta}).$$

#### k-th order Graph-Piecewise-Polynomial-Lasso

$$\hat{\beta}_k = \underset{\beta \in \mathbb{R}^p}{\text{argmin}} \quad \frac{1}{2N} \|y - X\beta\|_2^2 + \lambda \|\beta\|_1 + \lambda_k \|\Delta^{(k+1)}\beta\|_1$$

## Beyond linear models



#### Boosting algorithm 1

Input: dataset  $\{(X_i, Y_i, A_i)\}_{i=1}^n$ , number of iterations  $K$ , learning rate  $\eta$ , maximum of tree depth  $d$

- (1) Train  $\text{bst.plus1} = \text{XGBoost}(\{(X_i, Y_i); A_i = 1\}, K, \eta, d)$
- (2) Train  $\text{bst.minus1} = \text{XGBoost}(\{(X_i, Y_i); A_i = -1\}, K, \eta, d)$
- (3) Output the estimated optimal ITR:

$$\hat{D}(x) = \text{sign}(\text{bst.plus1}(x) - \text{bst.minus1}(x))$$

#### Boosting algorithm 2

Input: dataset  $\{(X_i, A_i, Y_i)\}_{i=1}^n$ , number of iterations  $K$ , shrinkage parameter  $\eta$  and maximum tree depth  $d$ .

- (1) Estimate the common effect  $\mu$ .
- (2) Train  $\text{bst} = \text{XGBoost}(\{(X_i, A_i \text{sign}(Y_i - \mu(X_i))\}, K, \eta, d)$  with weighted deviance loss
- (3) Output the estimated optimal ITR:

$$\hat{D}(x) = \text{sign}(\text{bst}(x))$$

# Takeaway

- Our work has revealed the importance of designing efficient statistical learning methods which adapt to the unique features of complex data
- This thesis leaves the door open for a more extensive investigation of efficient statistical learning for complex data
  - Other interesting types of complex data, e.g., high-frequency financial data, large network data
  - Unsupervised learning
- Our hope: develop more approaches which not only ensure statistical and computational soundness, but also provide easy-to-use, accessible software

- Our work has revealed the importance of designing efficient statistical learning methods which adapt to the unique features of complex data
- This thesis leaves the door open for a more extensive investigation of efficient statistical learning for complex data
  - Other interesting types of complex data, e.g., high-frequency financial data, large network data
  - Unsupervised learning
- Our hope: develop more approaches which not only ensure statistical and computational soundness, but also provide easy-to-use, accessible software

**Thank you!**